ALPhA Week 4

Jose Cruz and Andrew Hoyle

Summary of the last week

- Implemented Pix2Pix with our data
- Worked on building a CNN from the ground up to train from scratch.
- Found a typo

Pix2Pix

- Two networks competing (Generator and Discriminator)

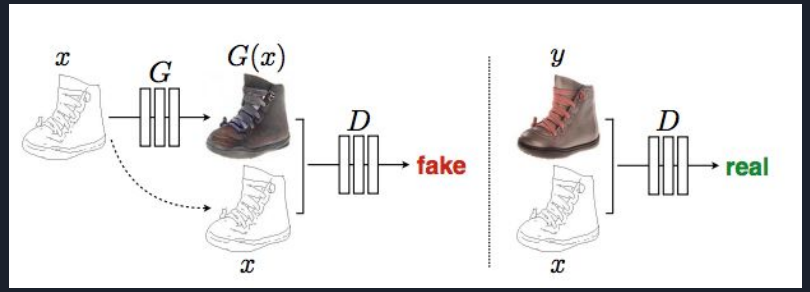

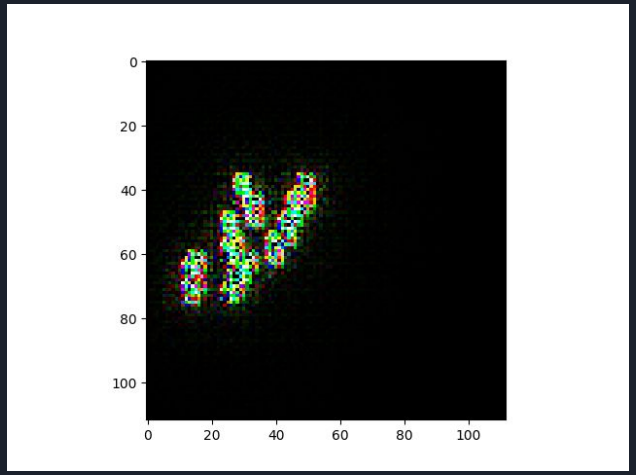

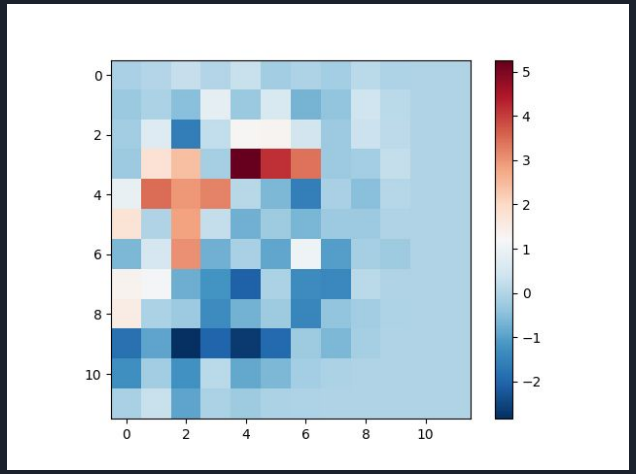

Pix2Pix Results

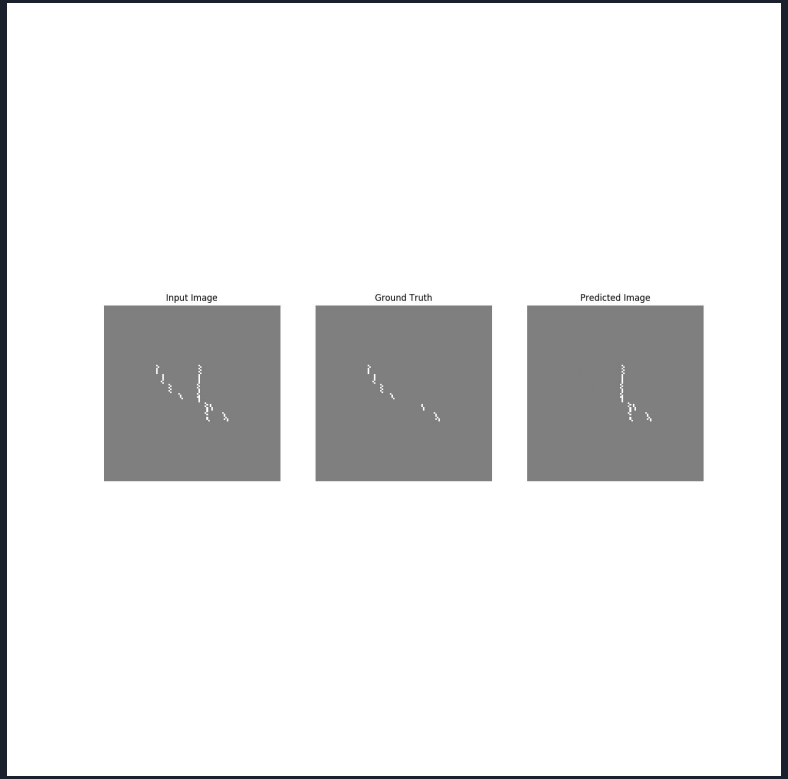

CNN from Scratch

- After finding that our results were around 50%, we decided to make a small scratch made CNN modeled after our VGG16 pretrained model
- The scratch made model would only have up to the second pooling layer

Layer (type)	Output Shape	Param #
conv2d (Conv2D)	(None, 112, 112, 64)	256
conv2d_1 (Conv2D)	(None, 112, 112, 64)	4160
max_pooling2d (MaxPooling2D)	(None, 56, 56, 64)	0
conv2d_2 (Conv2D)	(None, 56, 56, 128)	8320
conv2d_3 (Conv2D)	(None, 56, 56, 128)	16512
max_pooling2d_1 (MaxPooling2	(None, 28, 28, 128)	0
flatten (Flatten)	(None, 100352)	0
dense (Dense)	(None, 64)	6422592
activation (Activation)	(None, 64)	0
dense_1 (Dense)	(None, 1)	65
activation_1 (Activation)	(None, 1)	0

Total params: 6,451,905
Trainable params: 6,451,905
Non-trainable params: 0

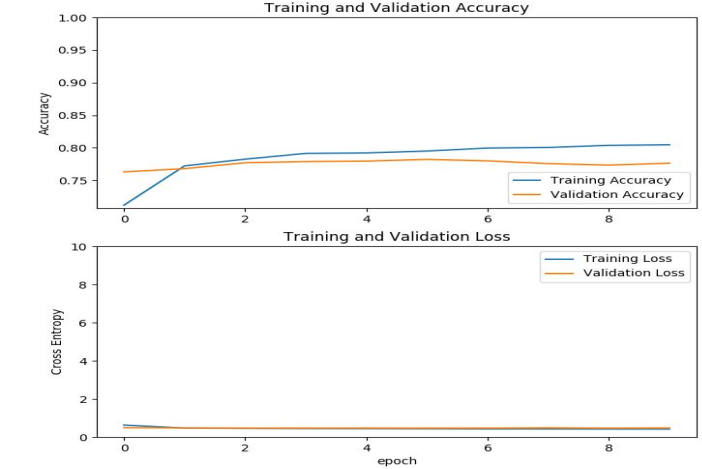

So about that typo...

- Accidentally gave the script the same dataset twice
- Fixed it
  - When training the entire model we reached 92% accuracy
  - Logistic regression accuracy: 89%

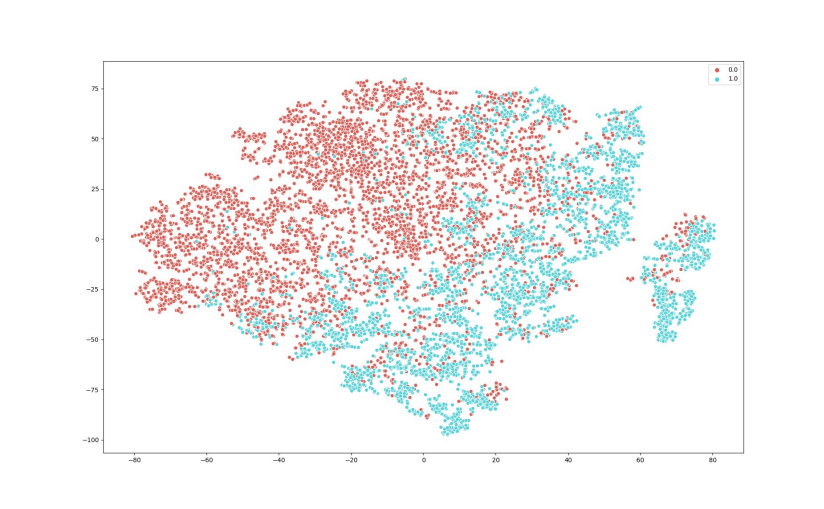

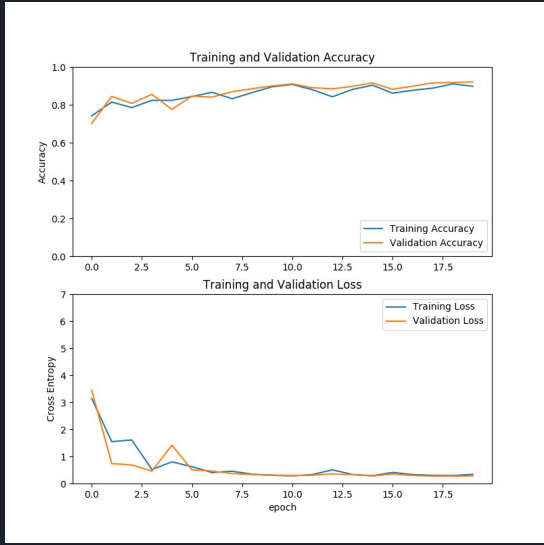

Typo Aftermath

- 92% accuracy is fun and all but Pix2Pix seems more applicable
  - No need to generate various projections of possible data

Goals for this week

- Meet with Gagik (discuss which method we prefer)
- Progress and fine tune Pix2Pix implementation
    - PatchGAN vs PixelGAN